

SpindleAI Application

•••

Michael Shomsky 4/2025

Requirements

- Agentic System with ≥ 2 agents
- Agents solve calculation problems
- Calculations utilize Math Tools
- When using python, bridge the gap and communicate to Rust
- Create virtual tools that generically solve problems
- When matching a virtual tool do use a local vector store

Requirements

- **Agentic System with ≥ 2 agents**
 - 3 agent system
- **Agents solve math problems**
 - Yes
- **Calculations utilize Math Tools**
 - MathToolbox organizes tools used by agents
- **When using python, bridge the gap and communicate to Rust**
 - MathToolbox organizes tools written in Rust and compiled to be used in Python
- **Create virtual tools that generically solve problems**
 - VirtualToolManager creates and manages virtual tools after successfully vetting a solution
- **When matching a virtual tool do use a local vector store**
 - Faiss is a library — developed by Facebook AI — that enables efficient similarity search used as a local vector store

MVP - 3 Agents + Virtual Tool Manager

3 Agents

- **Solver Agent** - Uses math tools and LLM to solve
- **CAS Agent** - Uses LLM to frame the problem to be solved by a Computer Algebra System
- **Verification Agent** - Uses LLM to solve

Tools

- **MathToolbox (Python)** sum, product, divide, subtract, power, sqrt, modulo, round_number
- **MathToolbox (Rust)** calculate_average
- **Virtual Tool Manager** - Manage functions

Model

- gpt-4o-mini
- temperature = 0

Agent Framework

- Langchain

Languages

- Python + Streamlit
- Rust

MVP - 2 Agents for discussion (Math_Planner -> Planner Execution)

3 Agents + 2 Agents for discussion

- Solver Agent - Uses math tools and LLM to solve
- CAS Agent - Uses LLM to frame the problem to be solved by a Computer Algebra System
- Verification Agent - Uses LLM to solve
- **Math Planner Agent**
- **Plan Execution Agent**

Tools

- **MathToolbox (Python)** sum, product, divide, subtract, power, sqrt, modulo, round_number
- **MathToolbox (Rust)** calculate_average
- **Virtual Tool Manager** - Manage functions

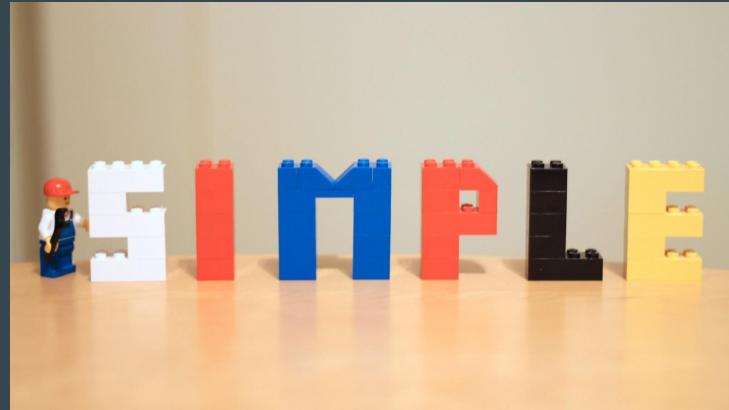
PLAN -> EXECUTE PATTERN

Plan: Create a plan on how to solve the problem and suggest tools to do so

Execute: Take the plan and use the available math toolkit resources to execute the plan

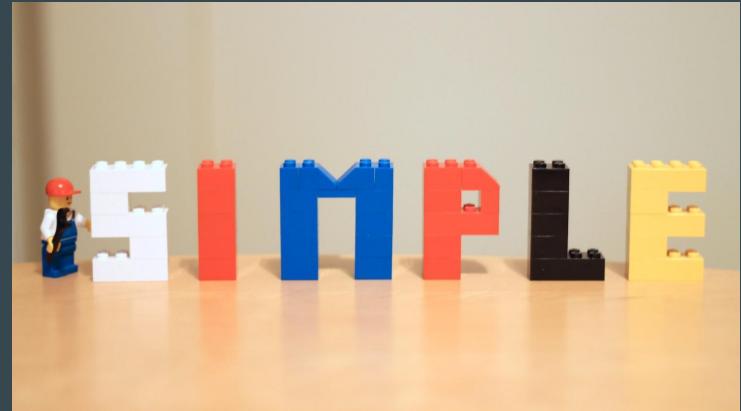
MVP - Keep it Simple

- Start with a simple 3 agents
 - Solver Agent
 - Uses MathToolbox's set of math functions
 - Uses LLM to guide an answer using MathToolbox
 - Verification Agent
 - Uses LLM to guide an thorough answer
 - CAS Solution Agent
 - Uses LLM to rephrase into a question consumable by a CAS
 - Uses a CAS to form an answer

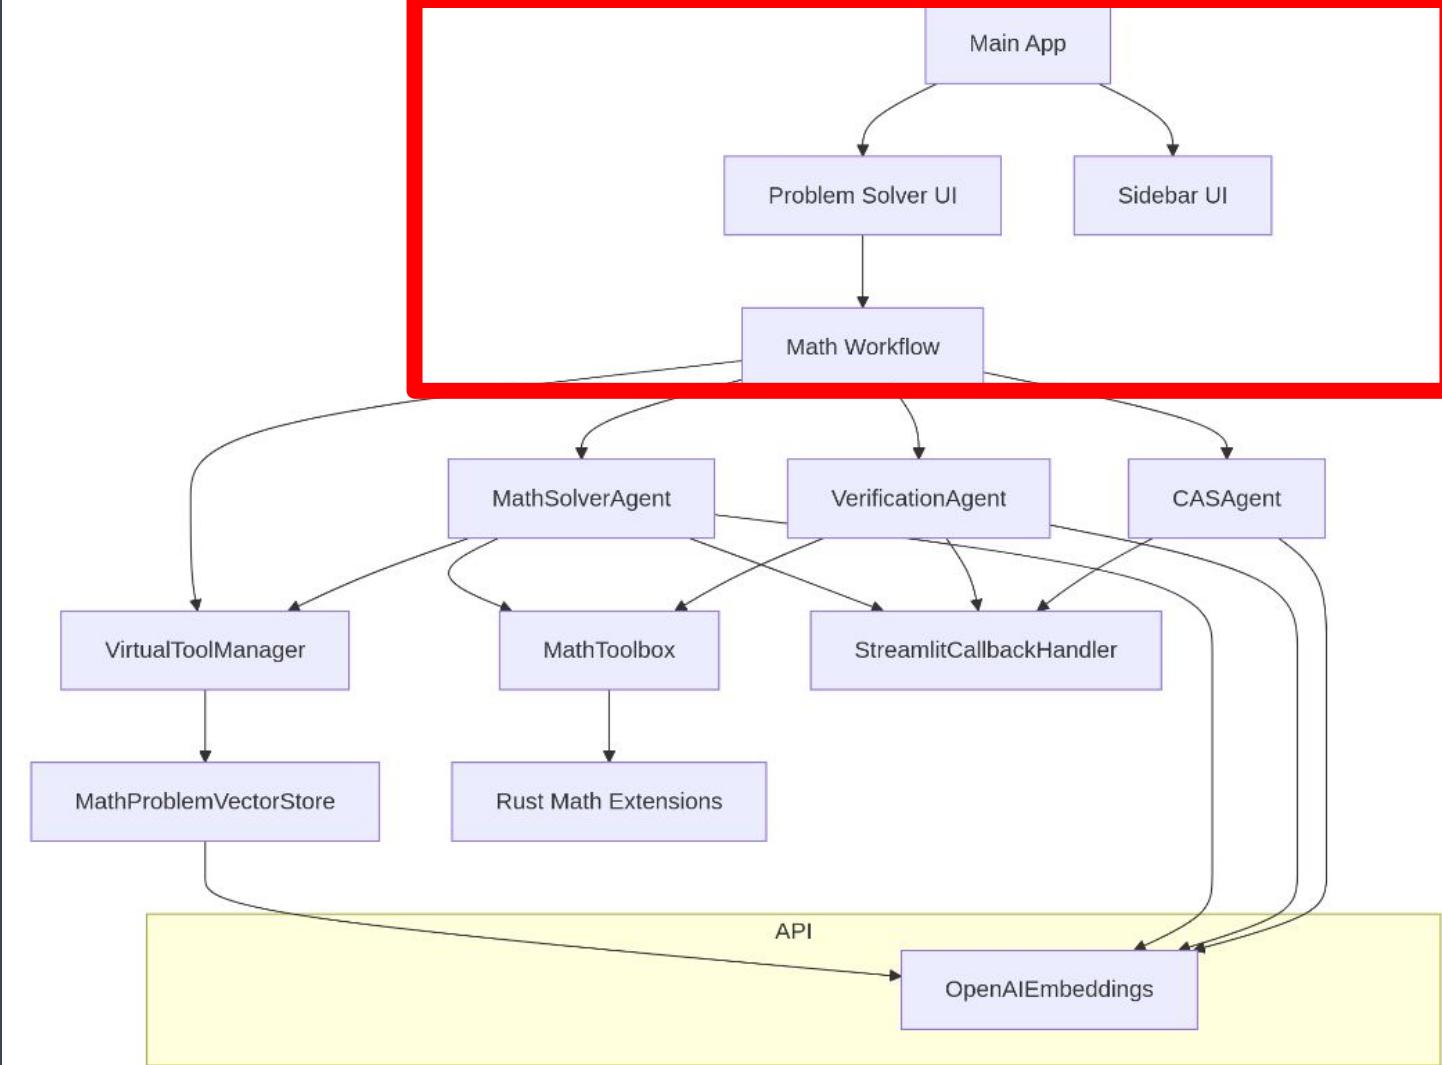


MVP - Majority Rules

- If any **two or three agents agree** on a solution, that solution is chosen as the majority solution.
- When there's no majority agreement:
 - If the **Solver and CAS agents agree** (but validation disagrees), we use their solution.
 - If the **CAS and Verification agents agree** (but solver disagrees), we use their solution.
 - Note: **Solver + Verification agent** is the the garden path to acceptance
- Only if we don't have any of the above cases, we fall back to the original behavior of using the solver solution with verification.



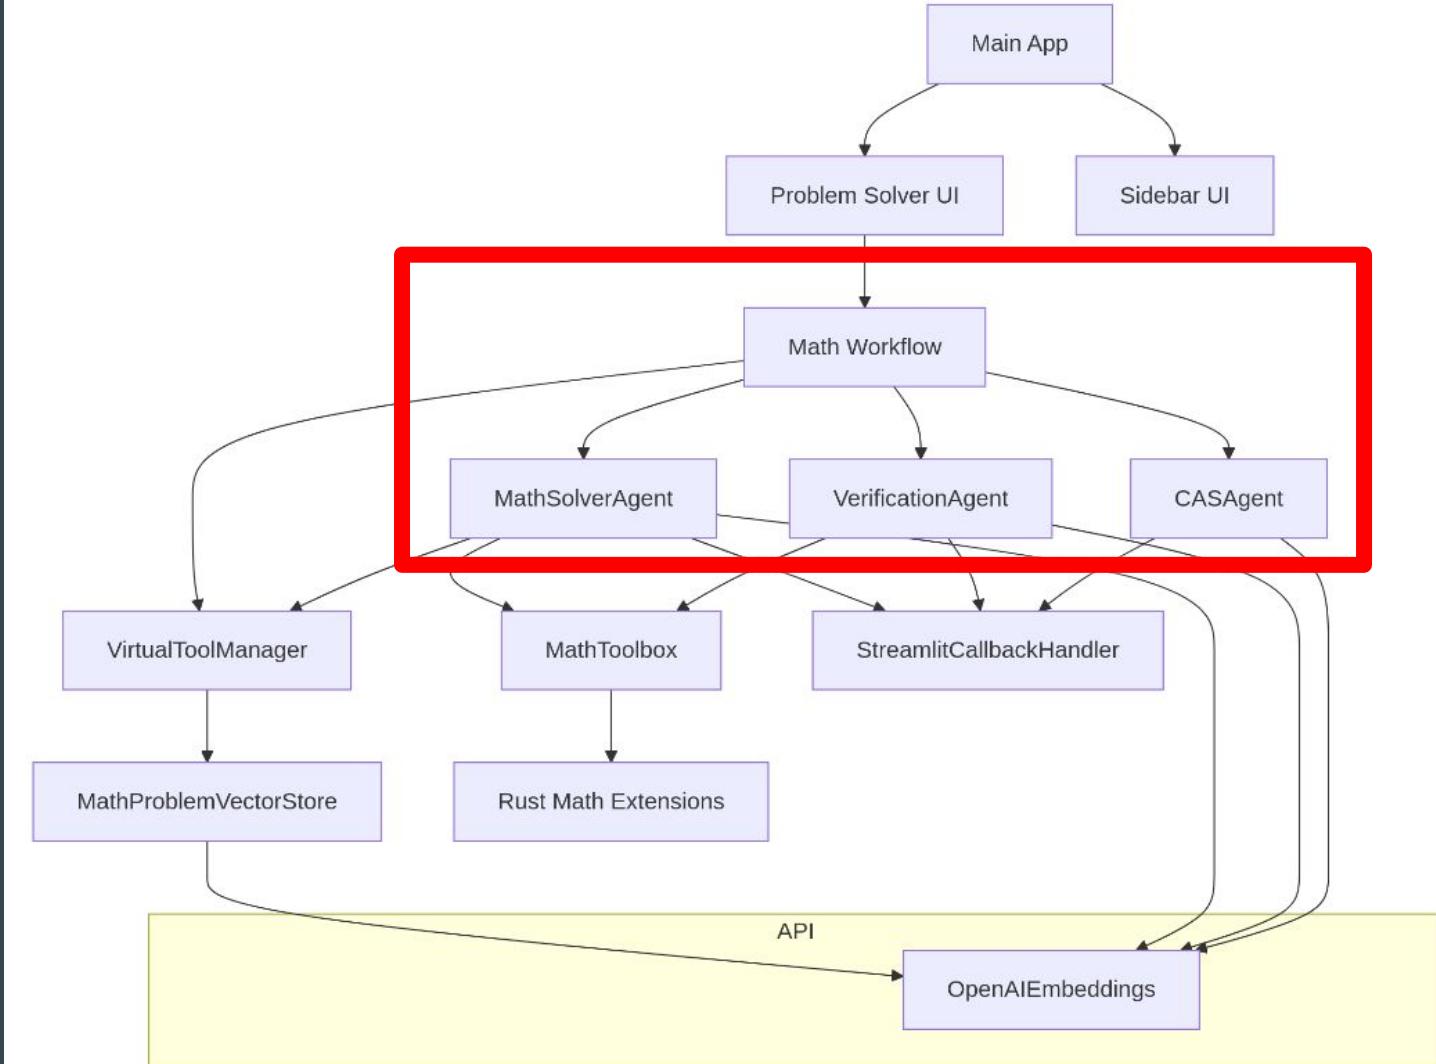
UI Drives a Math Workflow



Math Workflow Orchestrates 3 Agents

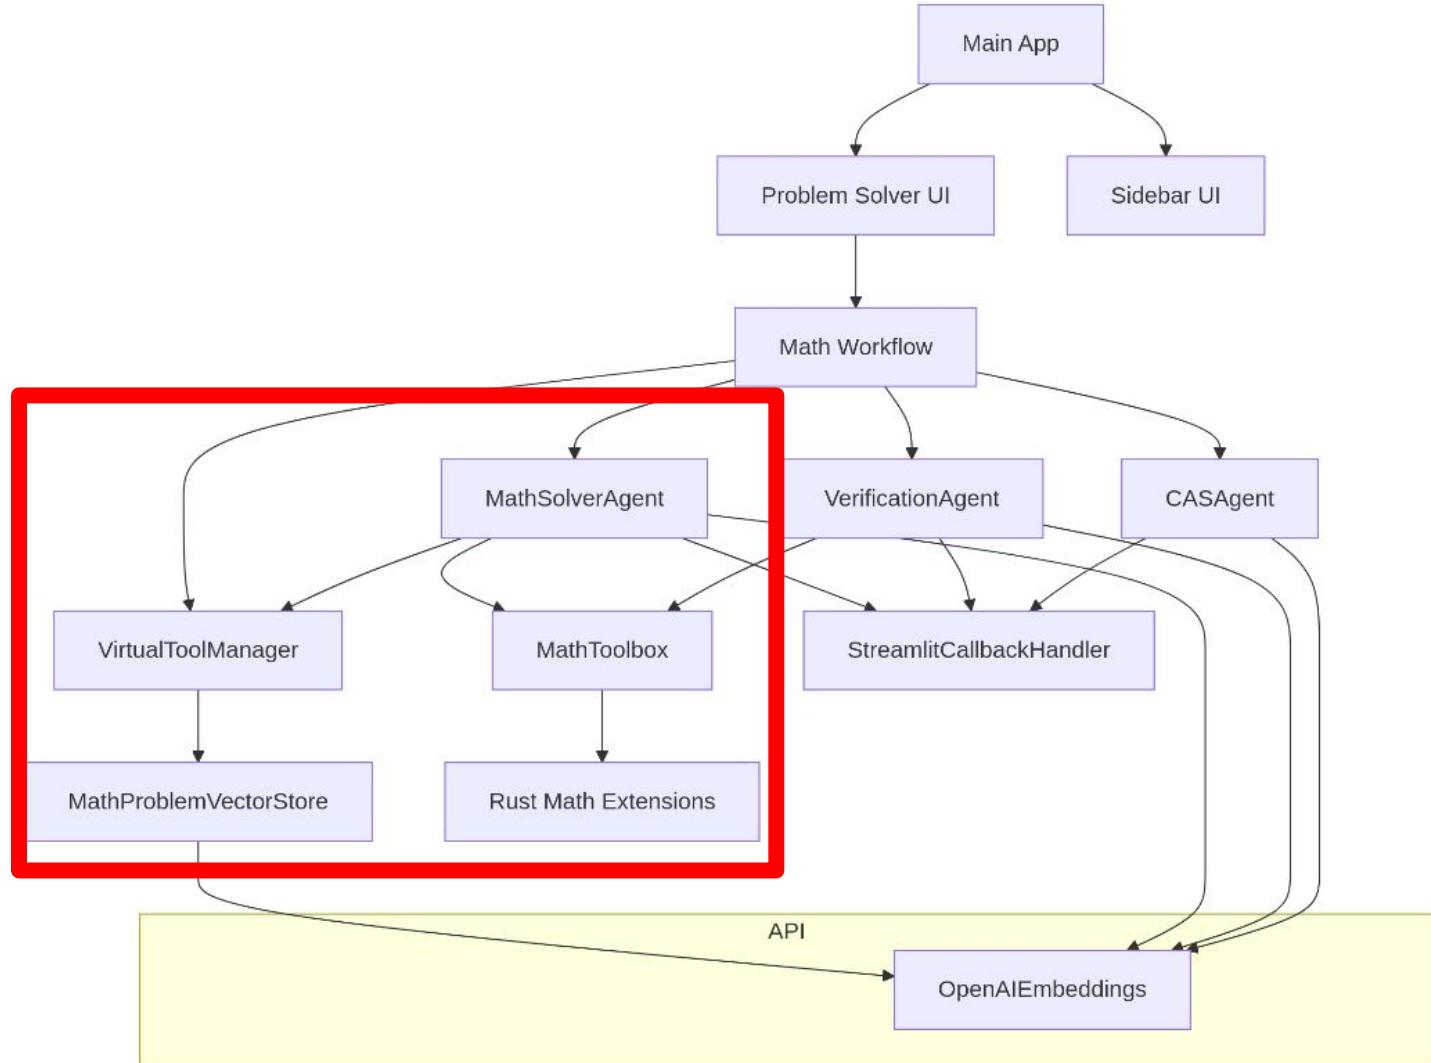
Or

Uses a Virtual Tool



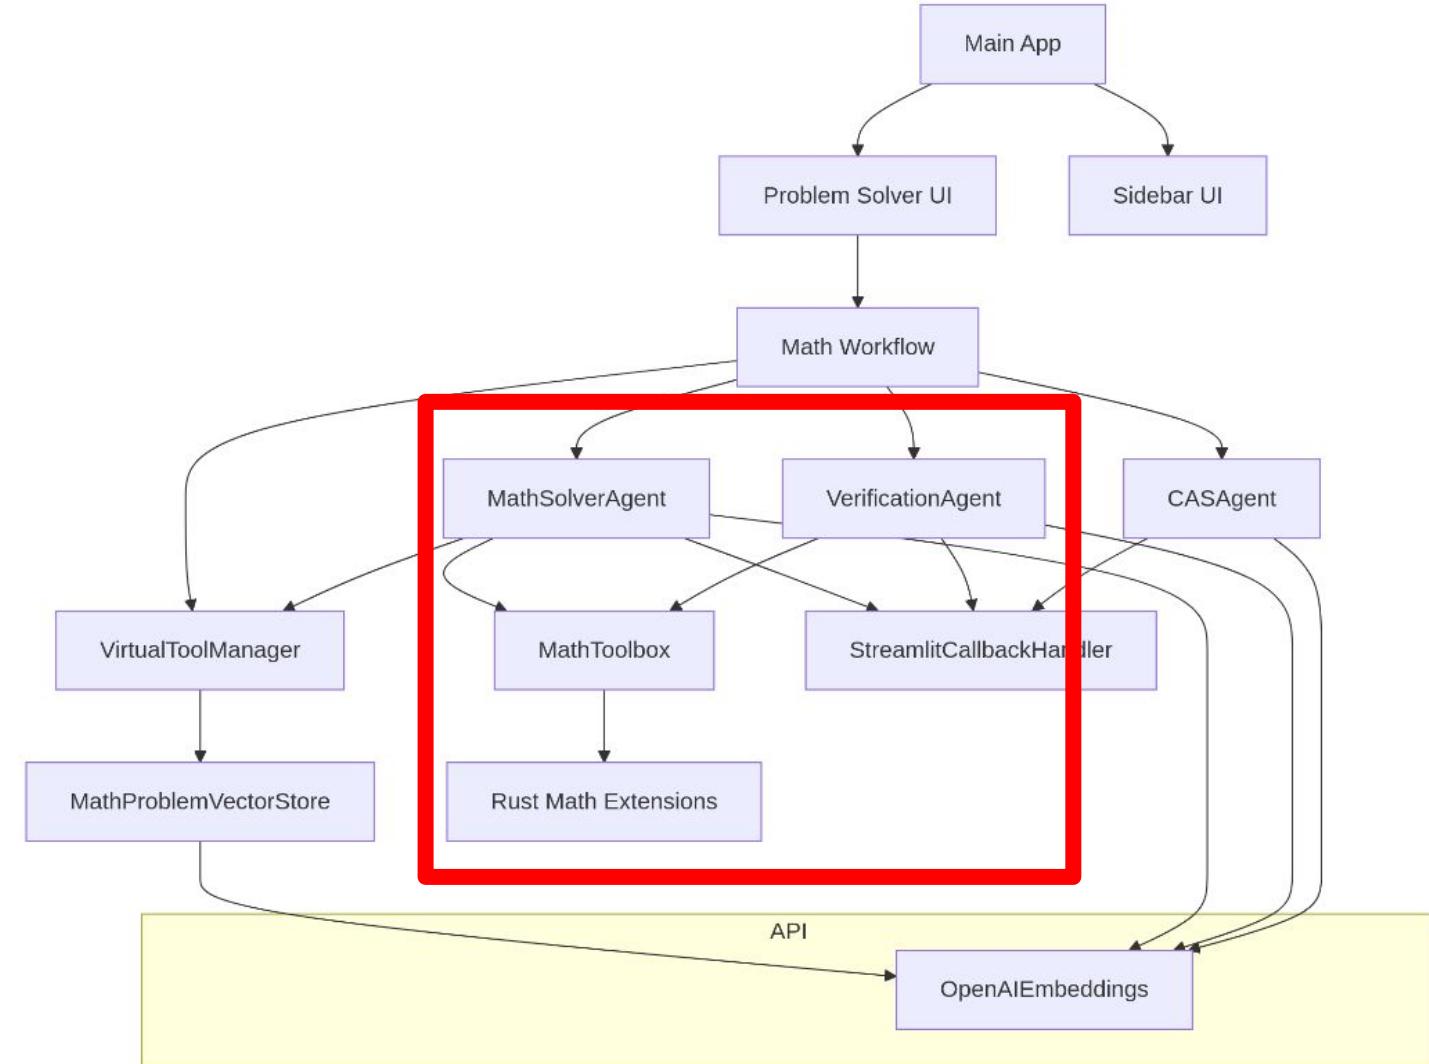
MathSolver Agents

uses Virtual Tool Manager to use an existing tool or record successful completion



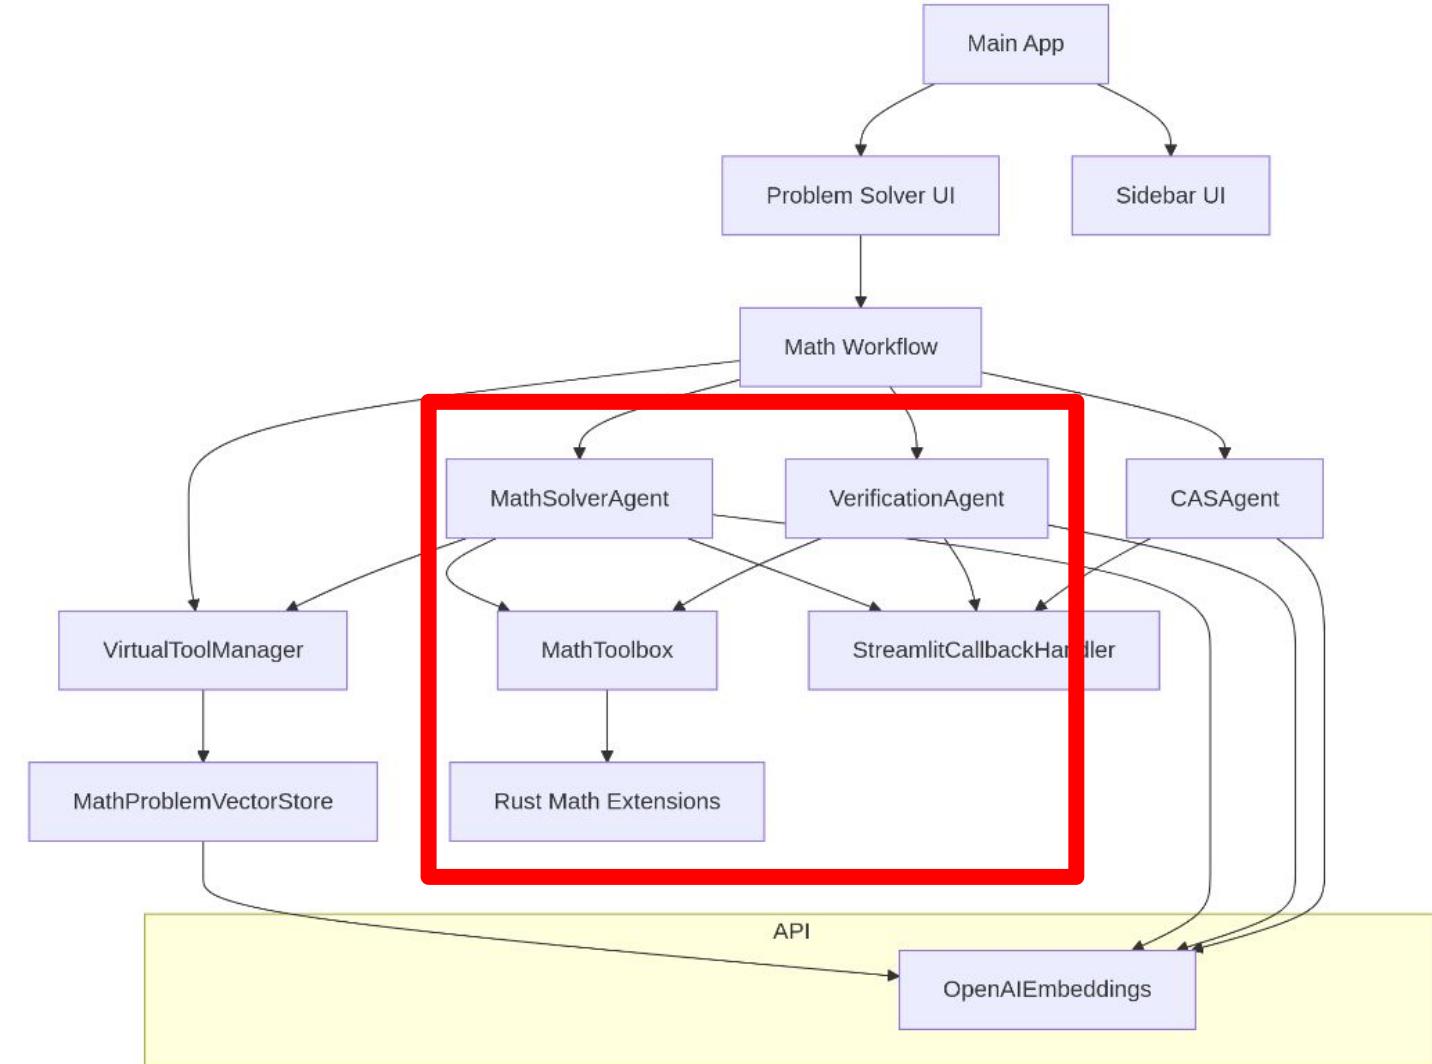
MathSolverAgent

Rely on
MathToolbox
For Math
Functions



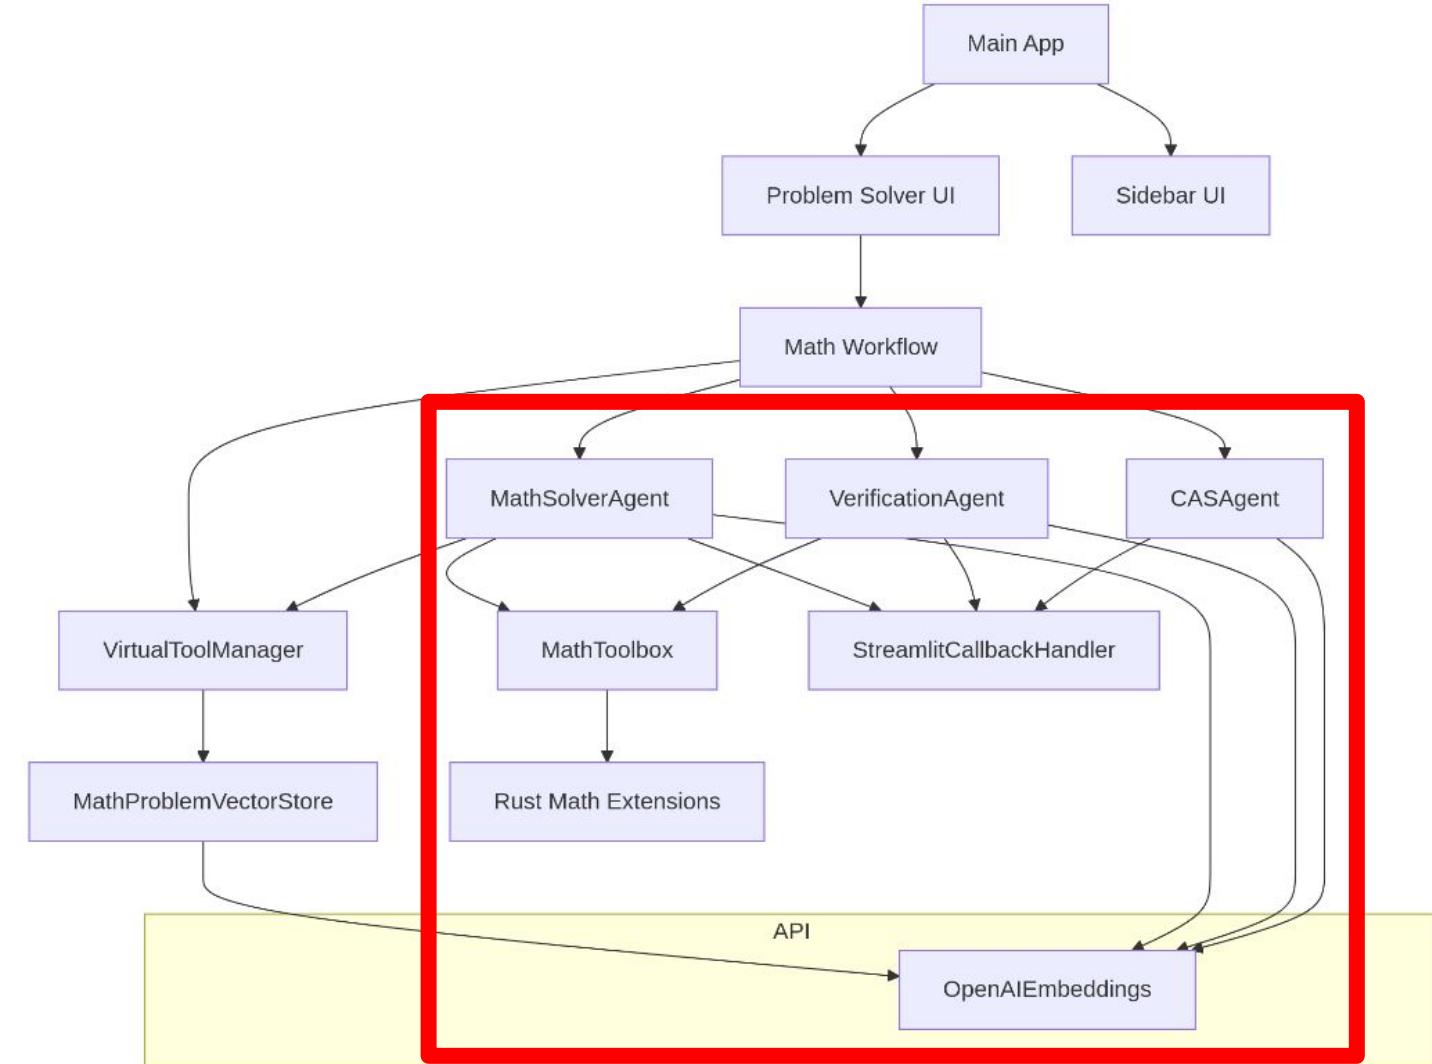
VerificationAgent

- LLM only for calculation
- (initially used math toolbox, but got deprecated)

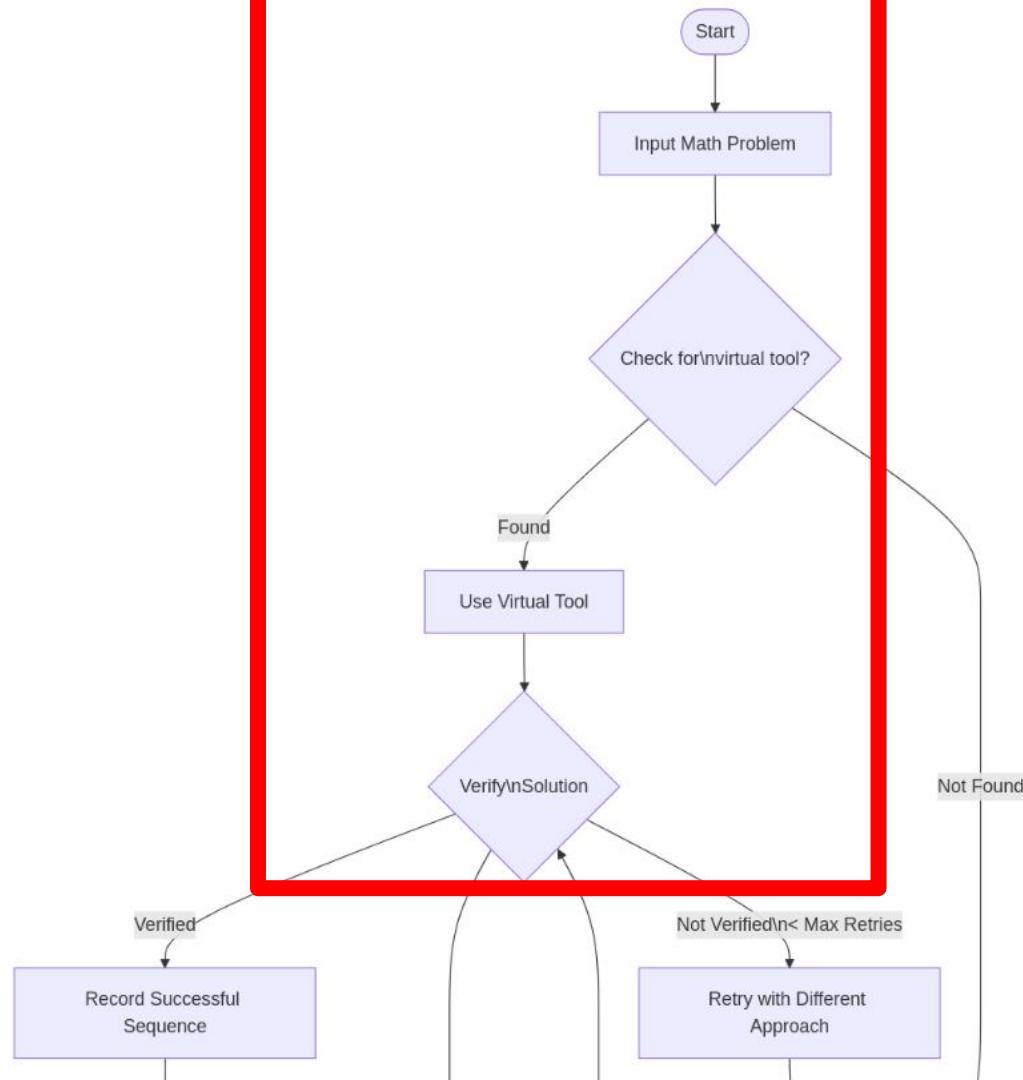


3 Agents

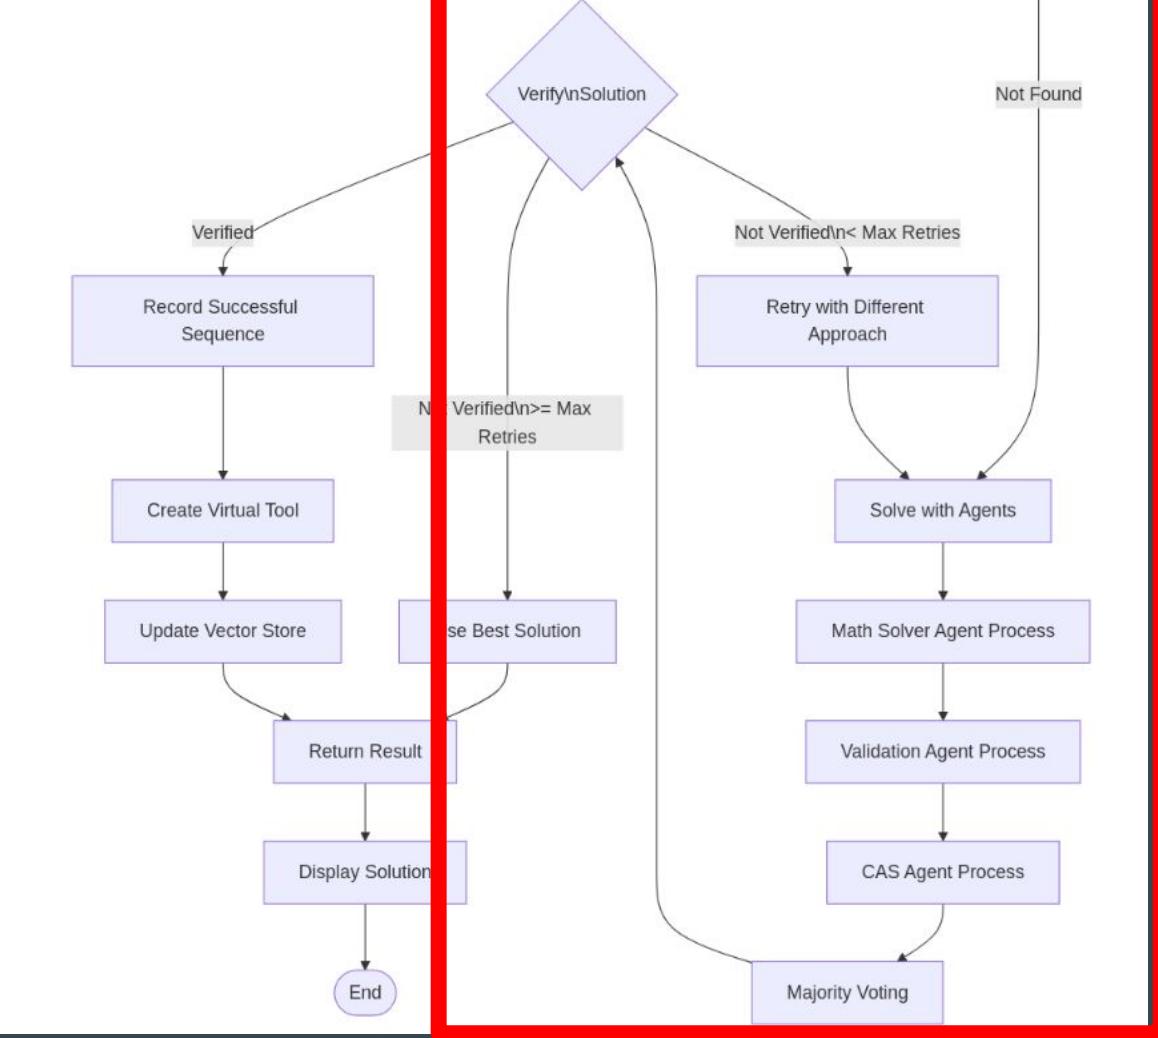
Rely on LLM integration
to work with
the math
question



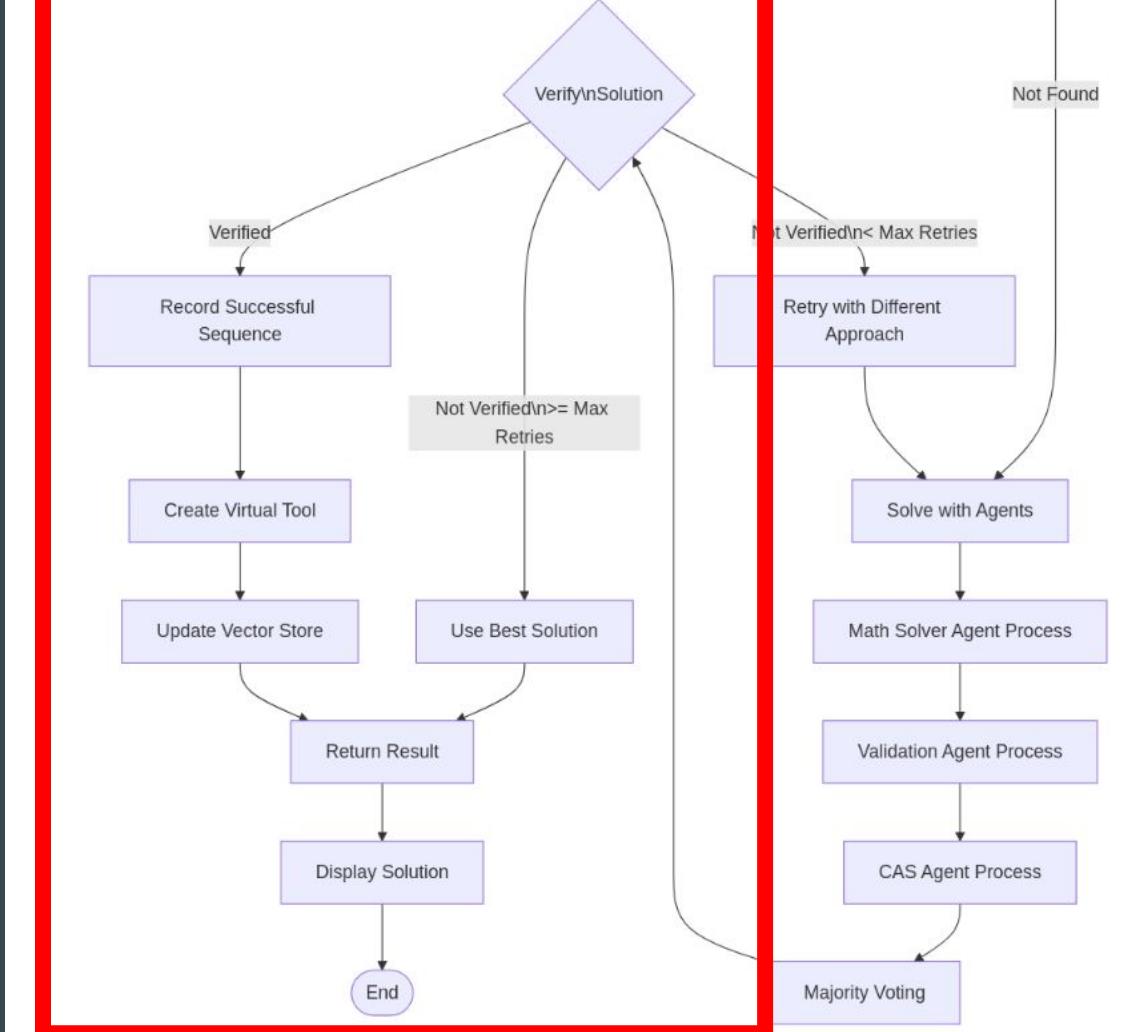
- Start by checking if there's a virtual tool.
- Use it if there is one
- Verify the Solution



- If verification fails
- Solve with Agents
- Use Majority Voting to verify answer



- Record successful Sequence
- Create a new virtual tool
- Update Vector Store
- Return/Display result



Demo - Python

Demo the python app using streamlit

- Live: <https://mvpmathsolver.replit.app/>
- Video: <https://youtu.be/Z3t02ZOIFfs>
- Local:
 - Branch (add-plan-exe)
 - docker compose up

The screenshot shows a Streamlit application with a dark theme. On the left, there's a sidebar titled 'Settings' with a toggle switch for 'Add Tool Errors'. Below it is a section titled 'Tool Statistics' showing call counts and percentages for various operations: sum (0 calls, 0.0%), product (2 calls, 0.0%), divide (0 calls, 0.0%), subtract (2 calls, 0.0%), power (2 calls, 0.0%), sqrt (0 calls, 0.0%), modulo (0 calls, 0.0%), and round_number (0 calls, 0.0%). At the bottom of this sidebar is a 'Virtual Tools' section listing 'VirtualTool_General_14dc969b: Solves General problems similar to: (2-3)*5^2'.

The main area has three columns. The left column is titled 'Enter a Math Problem' and contains a text input field with '(2-3)*5^2' and a dropdown menu showing the same expression. The middle column is titled 'Run Evaluation' and has a button labeled 'Run Evaluation on Sample Problems'. The right column is titled 'Sample Problems' and shows a text input field with '(2-3)*5^2' and a dropdown menu showing the same expression. Below these is a 'Selected problem:' dropdown with '(2-3)*5^2'. A 'Solve Problem' button is located at the bottom of this column. The bottom right column is titled 'Solution' and displays the result '-25'. The bottom left column is titled 'Verification' and contains a green box with a checked checkbox and the text 'The solution has been verified as correct! (Attempts: 1)'. The bottom right corner of the main area says 'VERIFIED: The proposed solution gives the correct'.

Demo

Demo the using the agents as a service locally

Start the API:

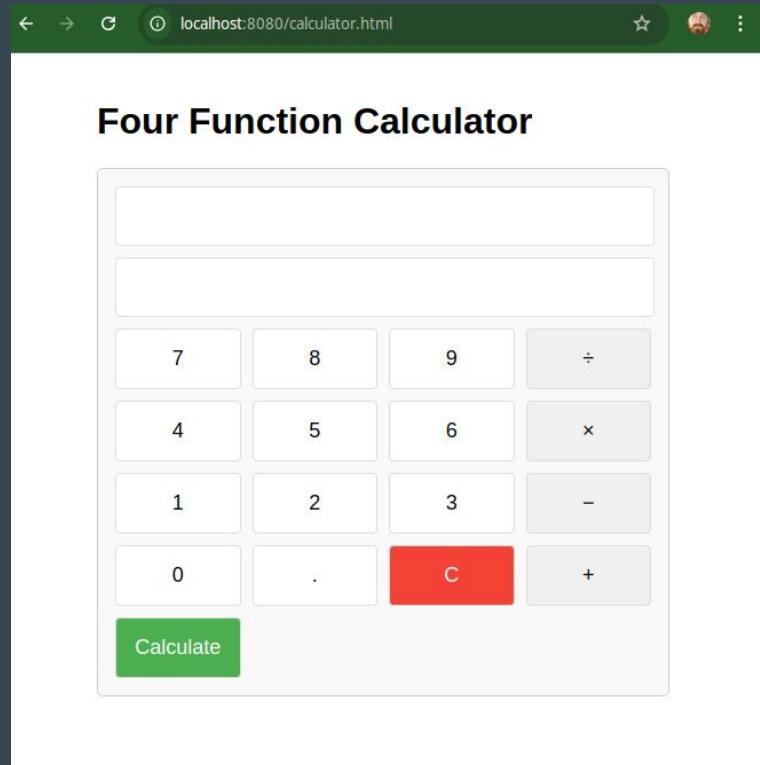
- Note: this is on the branch “add-plan-exe”
- Run the api (`./api/run.sh`)
- Run or load the page (`./api/serve_calculator.sh`)

Key details:

- Agents potentially as endpoints
- Shared instance of VirtualToolManager
- math_workflow on /api/tribunal for results

Demonstrates:

- Agent / Workflow Reusability
- Removes some of the streamlit clutter



Github

Invited <https://github.com/orgs/spindle-app/people>

https://github.com/mtshomskyieee/mvp_math_solver

- Python
- Rust

Thanks

- Next Steps
 - Scalable API using FastAPI + Kubernetes
 - Port the demo to Rust
- What I would do differently
 - Streamlit (and associated callbacks) made things more complicated. Next time I would go straight to API and cobble together a view
 - See 4 func calc
 - Caching is good, and makes repeat calls faster. `_workflow_cache` in `math_workflow` shouldn't be handled within the library. It's not that memoization is bad, so much as it is a maintenance problem. Creating a `ManagedCache` class might be a good first step, where you can add input:output pairs, clear it, turn it on/off would've been better.
 - I stuck with LangChain which now has LangGraph for workflows, and could be useful

Appendix

Additional slides

Appendix: Requirements

<https://spindle.notion.site/Coding-Project-Option-A-for-15757291437d804b87edf816a4212cdb>

This mission represents a stripped-down but realistic “toy version” of **the kind of multi-agent system Spindle AI is engineering** (including some actual challenges we’ve already faced):

1. **The Setup:** First, create ≥5 distinct, simple, deterministic tools that an LLM-based agent could call to help solve user-provided math problems (e.g. `SUM`, `DELTA`, `PRODUCT`, `QUOTIENT`, `MODULE`, `POWER`, `ABS`, `LOG`, `TRIG`, `SQRT`, `AVG`, `NODE`, `ROUND`, `UNION`, `INTERSECT`, `DIFFERENTIATE`, `INTEGRATE`, `FACTORIZIZE`, ... — the specific tools are entirely up to you).
 1. **Modify 1-2 of the most basic tools to intentionally but silently throw errors (and/or silently give incorrect answers) 30%-50% of the time the tool is called**. You may also want to include a basic `GET_USER_INPUT` tool for requesting input/clarification from a human user. (You can organize all tools in some form of “toolbox” if you want, but we’d prefer you do **not** hardcode a string listing all the tools, their docs, and their usage examples in a single prompt file or prompt mega-string anywhere in the project.)
 2. **The Architecture:** Prototype a multi-agent system with **at least 2 agents** and **at most 5 agents** (for whatever definition of “agent” you believe makes sense in this context), that discovers which tools are available and sequences tool calls to **reliably** solve basic user-provided math problems (or if you prefer, mathy word problems). The agents can **only** use the available tools (**including the unreliable tool[s]**) , i.e. no LLM-hallucinated arithmetic should be used for user-facing answers (even if that arithmetic is correct, as is increasingly the case among frontier models).
 1. You might well choose to include a lightweight planning, reasoning, and/or task decomposition layer in your prototype — but unless you have a compelling justification, all *user-facing* outputs (and most intermediate outputs) should be structured or semistructured, not unstructured.
 2. **Don’t hesitate to ask us for an OpenAI API key or Anthropic API key.** Otherwise, we’re happy to reimburse these costs after submission (*within reason/at Spindle’s discretion*).
 3. **The Twist:** When your prototype identifies a sequence of tool calls that reliably or fairly reliably solves a certain class of math problem(s) **based on successful execution(s)** , it should learn to do something like (e.g.) **memoize** or **semantically cache that sequence of tool calls as a single, idempotent new `VIRTUAL TOOL`** (i.e. some learning behavior akin to **“bundling” the tool calls into a *single* new idempotent tool, to which a *single* call can be made, which can be reliably invoked next time a math problem of the same or similar form is encountered).
 4. **The Finish Line:** Prove programmatically that your prototype works reasonably well (or at least that it could be *completed* to work reasonably well, if short on time).
 1. **Bonus points for using actual evals to show this.**
 1. **(If you’re an “evals-focused” candidate, consider reframing/approaching the entire task through the lens of an evals system instead, i.e. evals-driven development. Just tell us to judge your quality vs. emphasis vs. completion accordingly.)**
 5. **Bonus Points:**
 1. Create the math toolbox/interfaces in a non-Python language (ideally Rust, Go, or Typescript).
 2. If you decide to use a vector database anywhere, consider prototyping your own vector DB or VDB-like utility. (Not if this takes up all your time, though. It’s not the most important part.)
- **If you don’t have enough time for a project like this, or have alternate ideas, please let us know so we can find a path forward that we all feel good about!** Either way, we really look forward to seeing you through these next steps.

Agent Framework (LangChain)

- **LangChain-Based:** Structured, consistent agent implementation
- **Multiple Specialized Agents:** Solver, Verification, CAS, Planner, Executor
- **Model Selection:** gpt-4o-mini with temperature=0 for deterministic results
- **Tool Integration:** Streamlined interface between LLM and toolbox
- **Majority Voting:** Implements tribunal approach for consensus answers
- **Callback Design:** StreamlitCallbackHandler for real-time UI updates
- **Memory Management:** ConversationBufferMemory for stateful interactions
- **User Input Support:** Async interruption for gathering required information



MathToolbox

- **Core Operations:** sum, product, divide, subtract, power, sqrt, modulo, round_number
- **Reliability Control:** Tools can be set reliable/unreliable for testing
- **Error Simulation:** sum/product deliberately fail ~30-40% in unreliable mode
- **Complex Number Support:** Properly handles sqrt of negative numbers
- **Statistics Tracking:** Monitors usage and error rates per operation
- **Rust Integration:** Optimized avg function implemented in Rust
- **Testing:** Comprehensive test suite ensures correct math behavior

Rust Interoperability

- **High-Performance Extension:** Python calls native Rust function for avg operation
- **Fallback Mechanism:** Uses Python implementation if Rust unavailable
- **Dynamic Detection:** Auto-detects Rust extensions at runtime
- **Error Handling:** Graceful degradation if Rust function fails
- **Performance Benefits:** Faster calculations for repeated operations
- **Zero-Copy Interface:** Optimized data exchange between languages
- **Future-Ready:** Foundation for moving more compute-intensive operations to Rust

VirtualToolManager

- Creates reusable tools from successful math problem solutions
- Uses vector similarity to match new problems with existing tools
- Optimizes execution sequences & caches results for faster solutions
- Manages reliability through failure tracking & tool removal
- Provides pattern recognition for common math expressions
- Key metrics: 450+ lines of code with structural mapping algorithms

Walkthrough

1. **Problem Arrives:** Hash structure, normalize expression
2. **Tool Matching:** Seek similar structure via vector store lookup
3. **Execution:** Maps variables from new problem to original pattern
4. **Optimization:** Streamlines tool sequences (ex: sqrt of negative numbers)
5. **Tool Lifecycle:** Tracks performance, removes unreliable tools
6. **Vector Storage:** Employs embeddings for problem similarity
7. **Pattern Recognition:** Special handling for common forms $(a+b)/(c+d)$

Creating a New Math Agent: ex: MathTextSolutionAgent

Component	Explanation
Agent	math_text_solution_agent.py - Should be stored in /agents/ directory alongside the other agent files like cas_agent.py and solver_agent.py.
Adding to Workflow	In math_workflow.py, you would need to add the agent to the function parameters: <code>def math_workflow(..., math_text_solution_agent, ...)</code> [STREAMLIT SOLUTION ONLY] Because of the UI, you must initialize it in the <code>initialize_session_state()</code> function in main_app.py: if "math_text_solution_agent" not in st.session_state: st.session_state.math_text_solution_agent = MathTextSolutionAgent()
Adding to Tribunal In Workflow	To include the agent's results in the tribunal voting system, modify the voting logic in math_workflow.py by: 1. Adding the solution to the <code>normalized_solutions</code> dictionary: <code>normalized_solutions["text_solution"] = _normalize_solution(text_solution)</code> 2. Adding it to the "agent_solutions" in the result dictionary: <code>"agent_solutions": { ... "text_solution": text_solution, ... }</code>

Failsafe Mechanisms

When a Tool in the Solver Agent Fails

- **Error Detection and Reporting**
 - Tools track errors in `tool_stats` dictionary:

```
self.tool_stats["sum"]["errors"] += 1
```
 - All tools return formatted error messages:

```
return f"Error occurred: {str(e)}"
```
 - Example: In `math_toolbox.py`, division function explicitly checks for division by zero

When Solver Agent and Verification Agent Disagree

- **Retry Mechanism**
 - System attempts solution up to `MAX_VERIFICATION_RETRIES` (set to 5)
 - After disagreement, workflow tries alternative approaches

When Solver and CAS Agree but Verification Disagrees

- **Tribunal Voting System**
 - Implements a "majority rules" approach to resolve discrepancies:
 - CAS agent's symbolic verification becomes a third opinion to resolve deadlocks
 - If solver and CAS agree, their solution typically becomes the majority winner

Virtual Tool Deprecation

1. **Failure Counter System**
 - o Each virtual tool has an associated failure counter tracked in `tool_failure_counts` dictionary
 - o Failures are recorded via `record_tool_failure(problem_hash)` method
 - o Maximum allowed failures controlled by `max_failures` parameter (default: 3)
2. **Verification and Removal Process**
 - o When a virtual function fails verification or raises an exception during execution
 - o The failure counter is incremented: `self.tool_failure_counts[problem_hash] += 1`
 - o When threshold is reached: `if self.tool_failure_counts[problem_hash] >= self.max_failures:`

Tool is deleted from multiple locations:

```
del self.virtual_tools[problem_hash]          # Remove

functionself.vector_store.remove_problem(problem_hash)  # Remove from vector store

del self.successful_sequences[problem_hash]  # Prevent recreation
```

3. **Implementation in Workflow**
 - o Used in `math_workflow.py` to manage unreliable virtual tools
 - o If a virtual tool fails verification, the system falls back to standard solver
 - o Logs notification: "Virtual tool has been removed due to reaching max failures"

This mechanism ensures unreliable virtual functions are automatically purged, maintaining solution quality while allowing the system to learn and create new, more reliable tools over time.