

# SpindleAI Application

•••

Michael Shomsky 3/2025

# MVP - 3 Agents + Virtual Tool Manager

## 3 Agents

- **Solver Agent** - Uses math tools and LLM to solve
- **CAS Agent** - Uses LLM to frame the problem to be solved by a Computer Algebra System
- **Verification Agent** - Uses LLM to solve

## Tools

- **MathToolbox (Python)** sum, product, divide, subtract, power, sqrt, modulo, round\_number
- **MathToolbox (Rust)** calculate\_average
- **Virtual Tool Manager** - Manage functions

## Model

- gpt-4o-mini
- temperature = 0

## Agent Framework

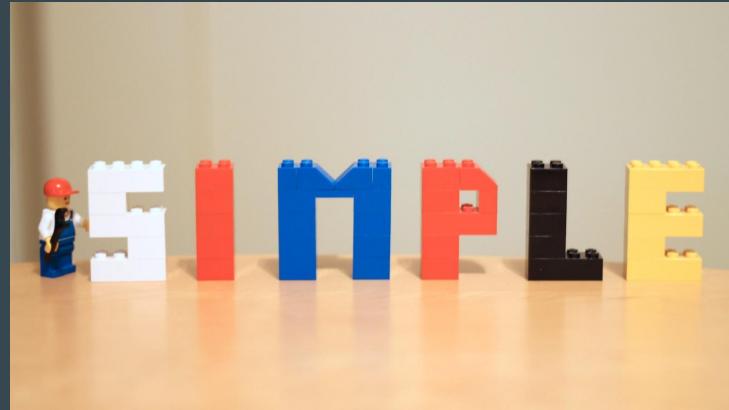
- Langchain

## Languages

- Python + Streamlit
- Rust

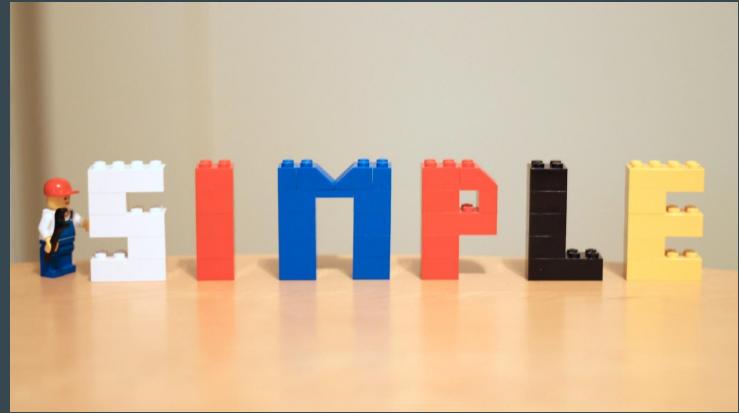
# MVP - Keep it Simple

- Start with a simple 3 agents
  - Solver Agent
    - Uses MathToolbox's set of math functions
    - Uses LLM to guide an answer using MathToolbox
  - Verification Agent
    - Uses LLM to guide an thorough answer
  - CAS Solution Agent
    - Uses LLM to rephrase into a question consumable by a CAS
    - Uses a CAS to form an answer

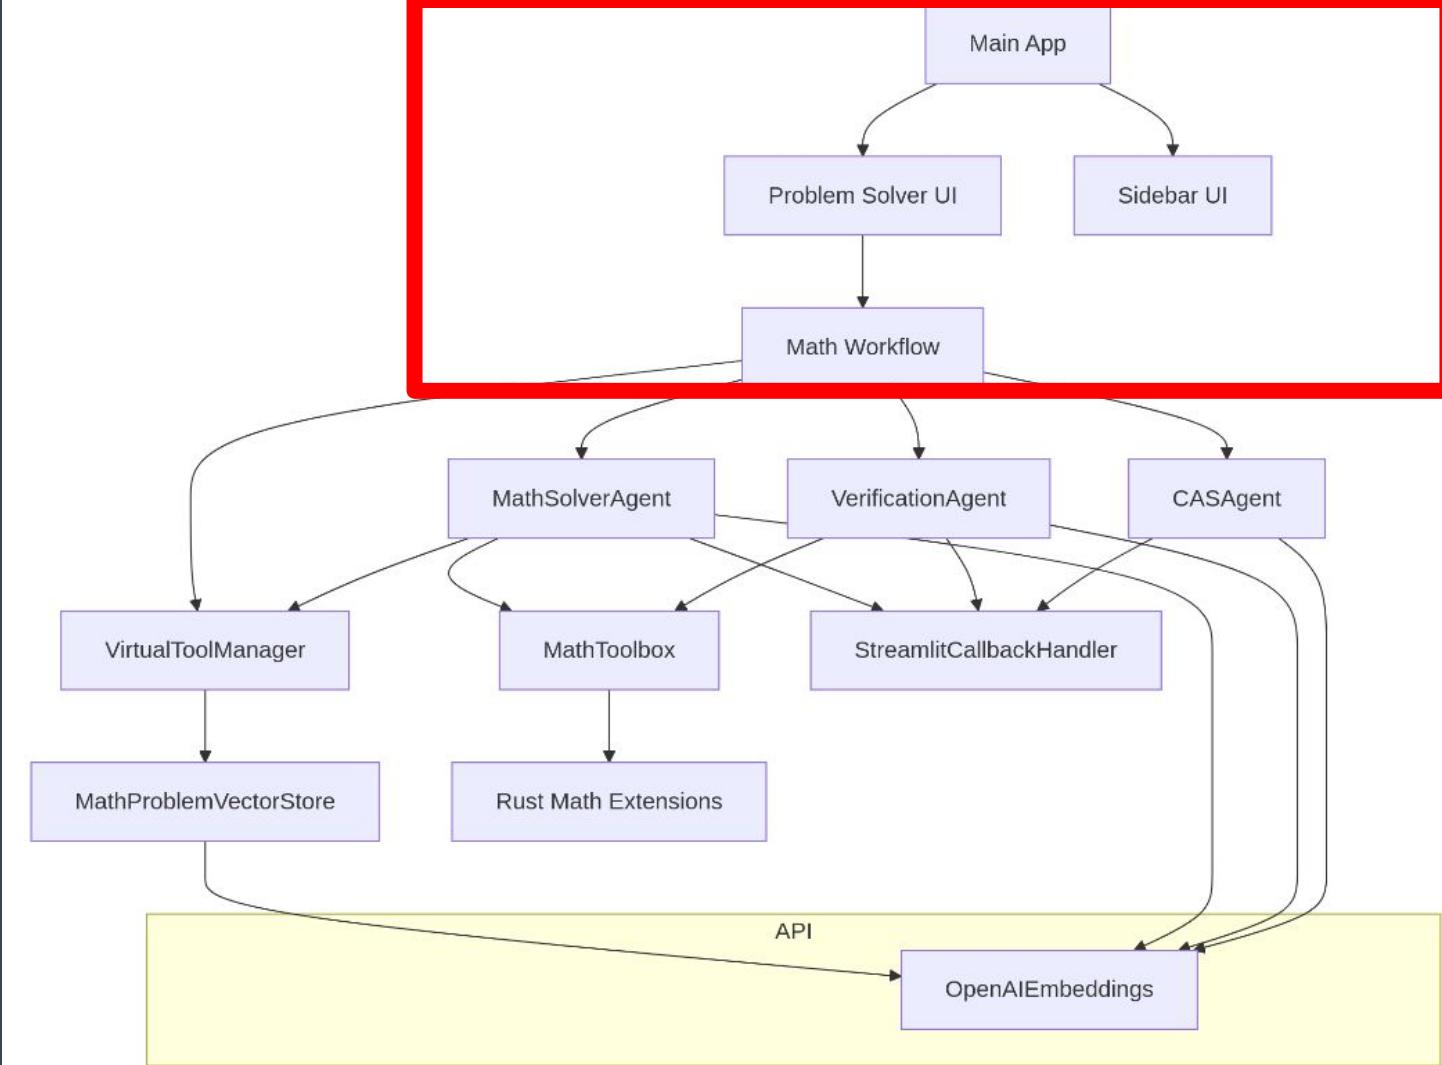


# MVP - Majority Rules

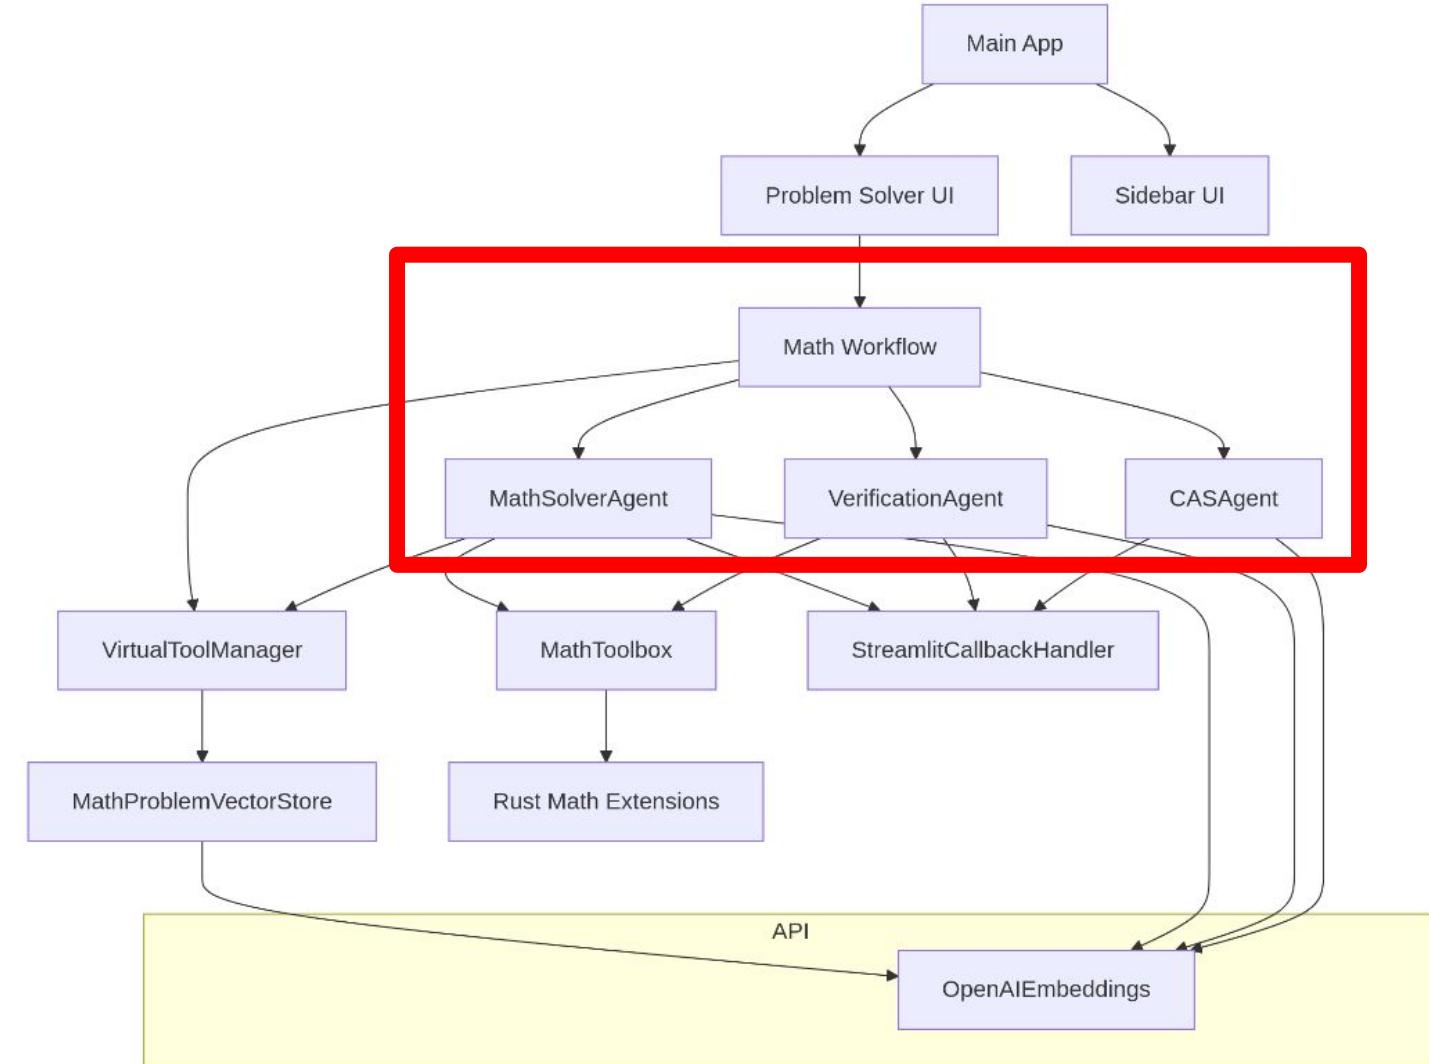
- 1. If any two or three agents agree on a solution, that solution is chosen as the majority solution.
- 2. If there's no majority agreement, we check two specific cases:
  - If the solver and CAS agents agree (but validation disagrees), we use their solution.
  - If the CAS and validation agents agree (but solver disagrees), we use their solution.
- 3. Only if we don't have any of the above cases, we fall back to the original behavior of using the solver solution with verification.



# UI Drives a Math Workflow

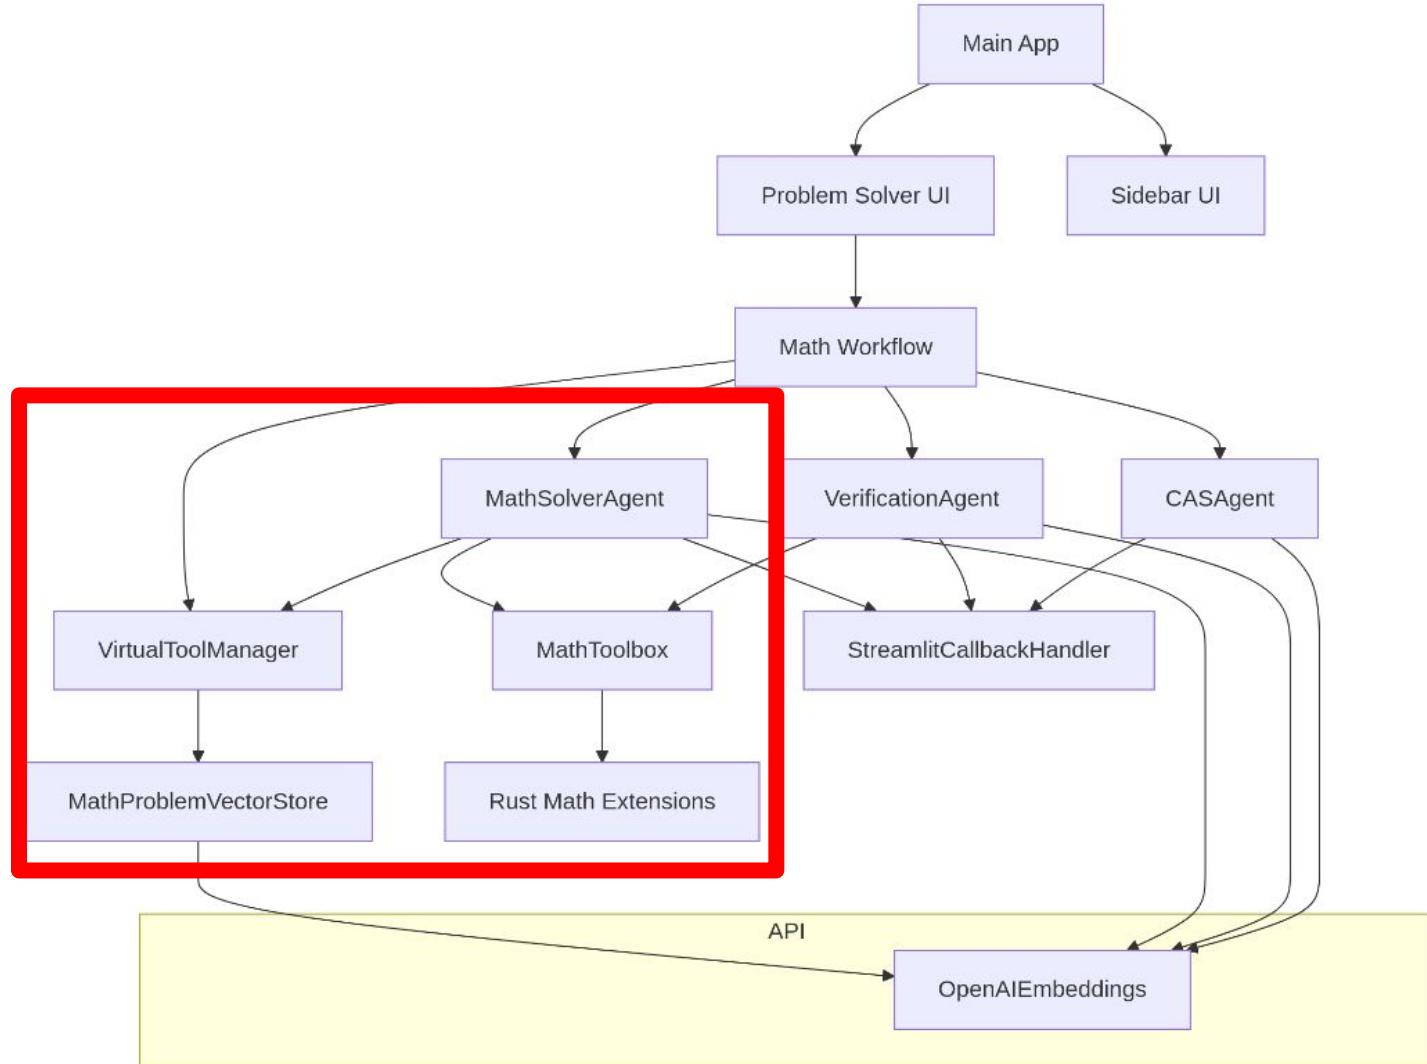


# Math Workflow Orchestrates 3 Agents



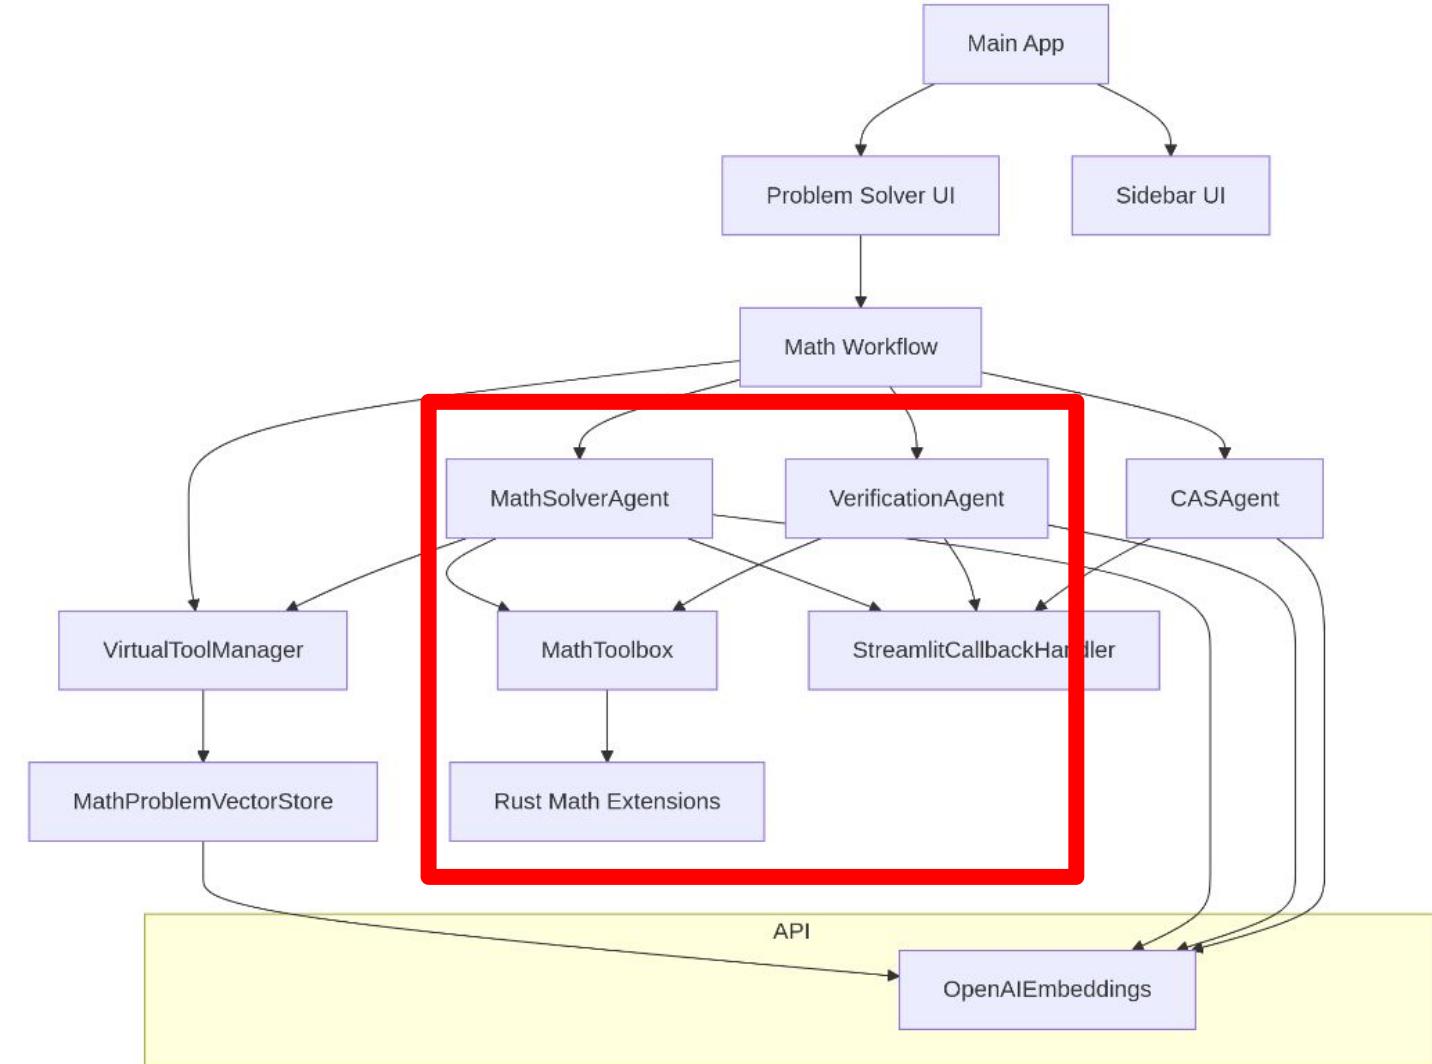
# MathSolver Agents

uses Virtual Tool Manager to get relevant function



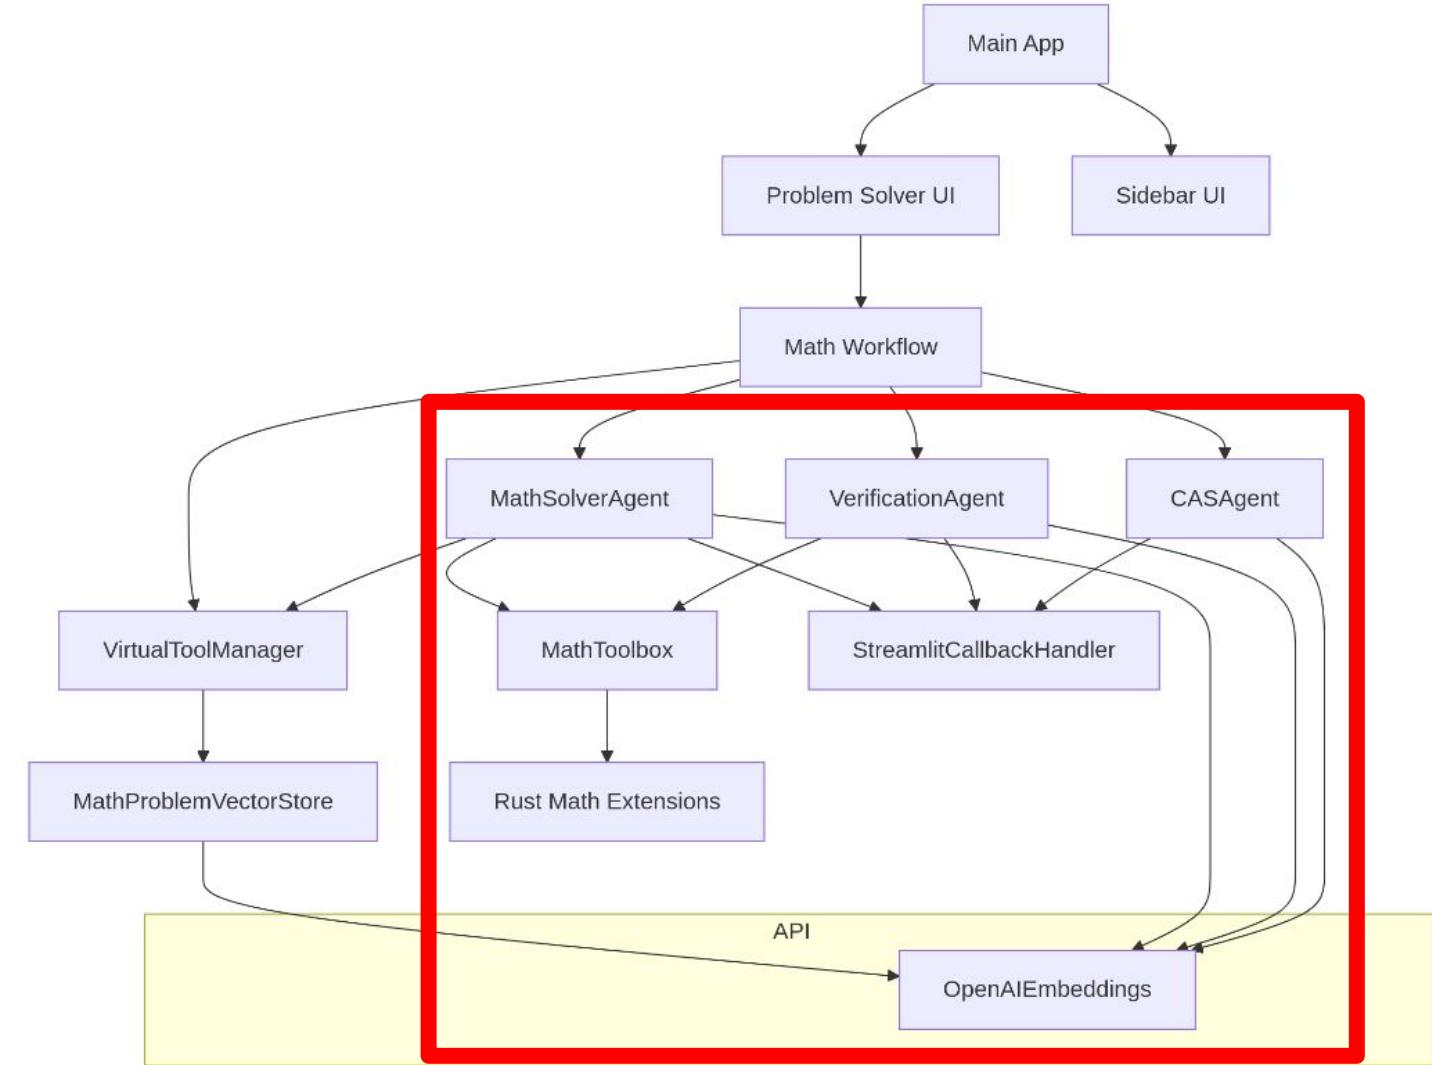
# 2 Agents

Rely on  
MathToolbox  
For Math  
Functions

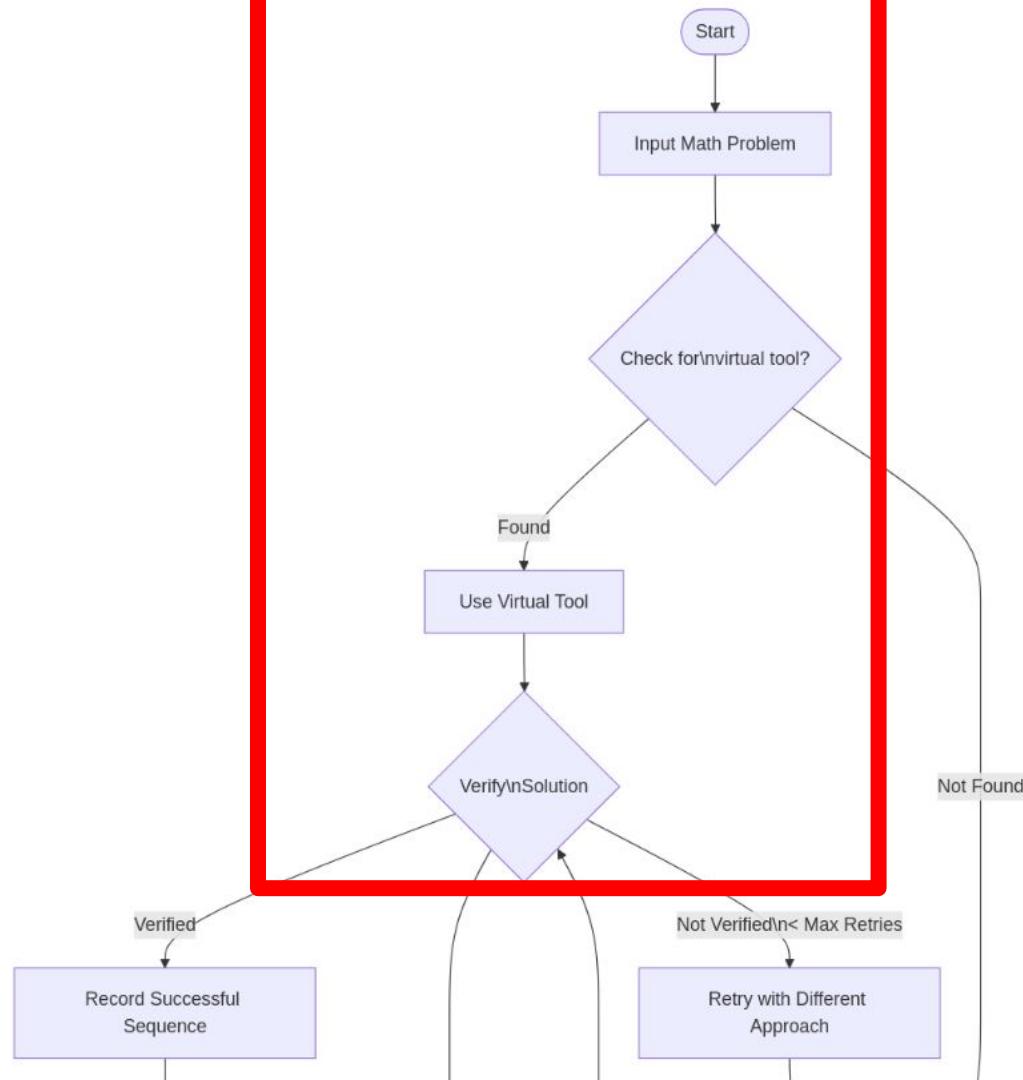


# 3 Agents

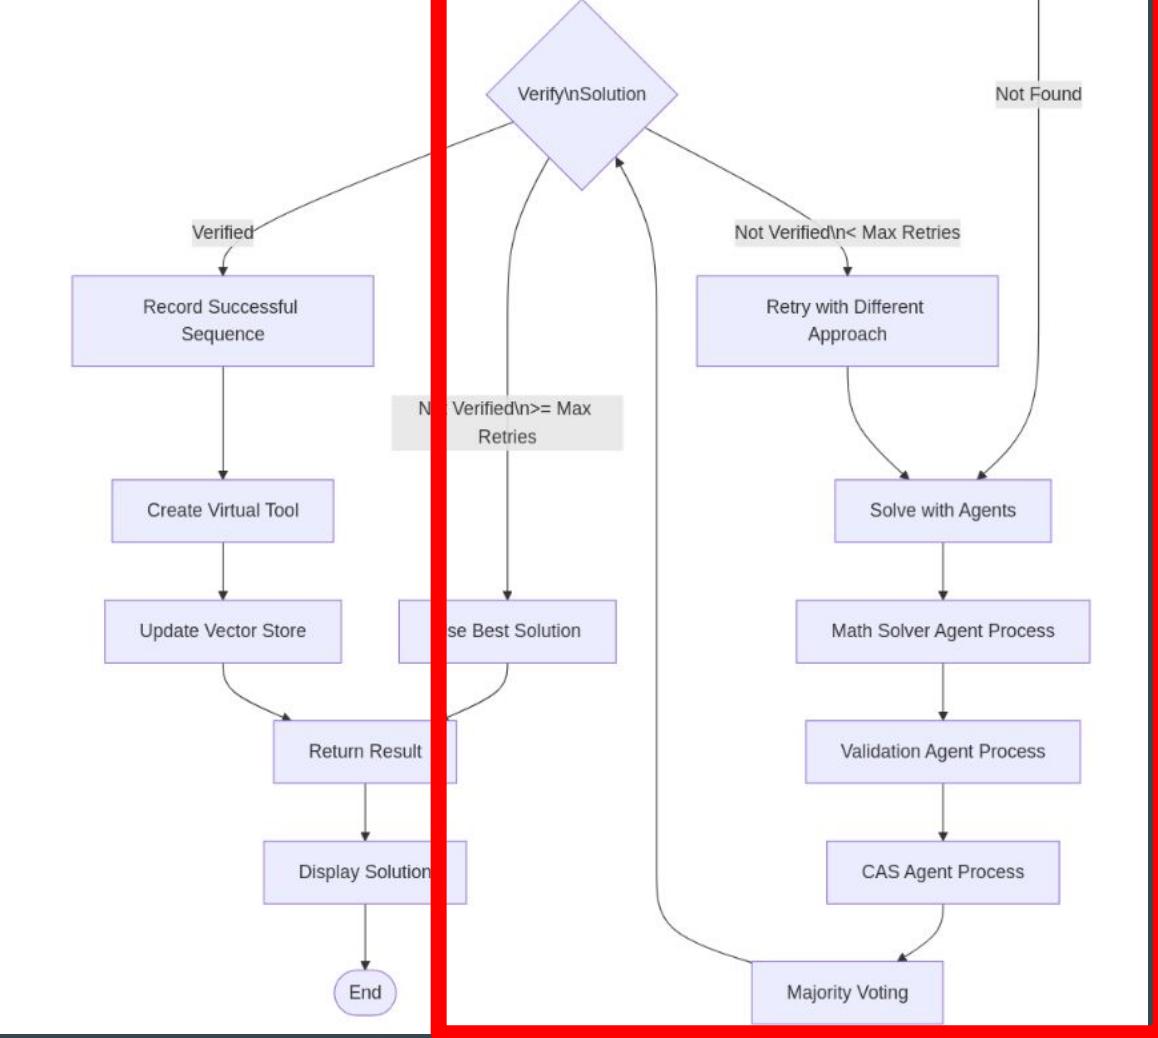
Rely on LLM integration  
to work with  
the math  
question



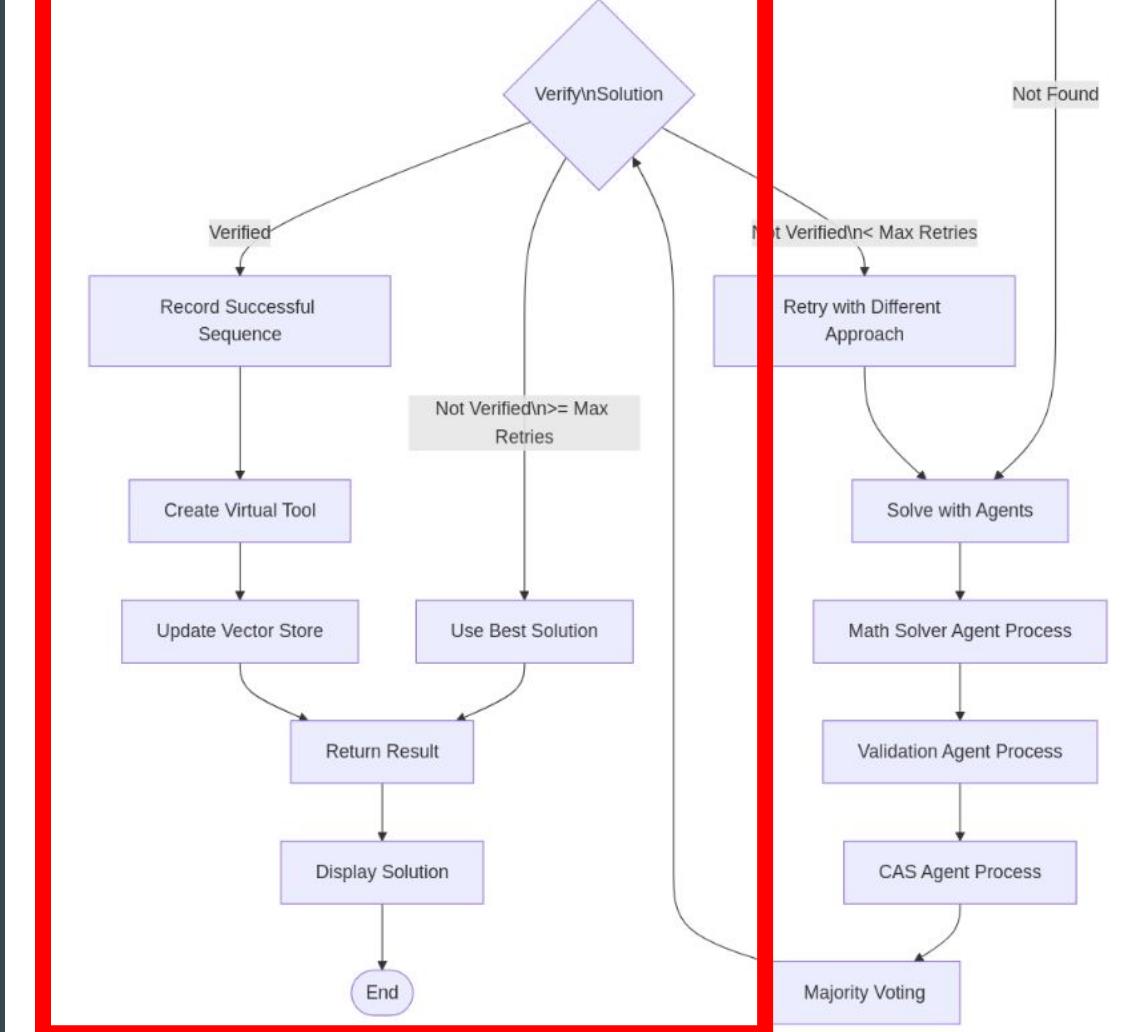
- Start by checking if there's a virtual tool.
- Use it if there is one
- Verify the Solution



- If verification fails
- Solve with Agents
- Use Majority Voting to verify answer



- Record successful Sequence
- Create a new virtual tool
- Update Vector Store
- Return/Display result



# Demo - Python

Demo the python app using streamlit

- Live: <https://mvpmathsolver.replit.app/>
- Video: <https://youtu.be/Z3t02ZOIFfs>

The screenshot shows the MVP Maths Solver application interface. It has a dark theme with light-colored text and buttons.

**Settings**: A toggle switch labeled "Add Tool Errors" is turned off.

**Tool Statistics**:

Tool	Calls	Percentage
sum	0 calls	↑ 0.0%
product	2 calls	↑ 0.0%
divide	0 calls	↑ 0.0%
subtract	2 calls	↑ 0.0%
power	2 calls	↑ 0.0%
sqrt	0 calls	↑ 0.0%
modulo	0 calls	↑ 0.0%
round_number	0 calls	↑ 0.0%

**Enter a Math Problem**: A text input field contains the expression  $(2-3)^5 \cdot 2$ . Below it is a dropdown menu showing the same expression.

**Run Evaluation**: A button labeled "Run Evaluation on Sample Problems".

**Sample Problems**: A dropdown menu also displays the expression  $(2-3)^5 \cdot 2$ .

**Solution**: The result "-25" is shown.

**Verification**: A green box indicates "The solution has been verified as correct! (Attempts: 1)".

**Sequence**: A note at the bottom states "Sequence: subtract → power → product".

# Github

Invited <https://github.com/orgs/spindle-app/people>

[https://github.com/mtshomskyieee/mvp\\_math\\_solver](https://github.com/mtshomskyieee/mvp_math_solver)

- Python
- Rust (ported from Python)

# Thanks

- Next Steps
  - Scalable API using FastAPI + Kubernetes
  - Port the demo to Rust

# Appendix: Requirements

<https://spindle.notion.site/Coding-Project-Option-A-for-15757291437d804b87edf816a4212cdb>

This mission represents a stripped-down but realistic “toy version” of **the kind of multi-agent system Spindle AI is engineering** (including some actual challenges we’ve already faced):

1. **The Setup:** First, create ≥5 distinct, simple, deterministic tools that an LLM-based agent could call to help solve user-provided math problems (e.g. `SUM`, `DELTA`, `PRODUCT`, `QUOTIENT`, `MODULE`, `POWER`, `ABS`, `LOG`, `TRIG`, `SQRT`, `AVG`, `NODE`, `ROUND`, `UNION`, `INTERSECT`, `DIFFERENTIATE`, `INTEGRATE`, `FACTORIZIZE`, ... — the specific tools are entirely up to you).
    1. **Modify 1-2 of the most basic tools to intentionally but silently throw errors (and/or silently give incorrect answers) 30%-50% of the time the tool is called**. You may also want to include a basic `GET_USER_INPUT` tool for requesting input/clarification from a human user. (You can organize all tools in some form of “toolbox” if you want, but we’d prefer you do **not** hardcode a string listing all the tools, their docs, and their usage examples in a single prompt file or prompt mega-string anywhere in the project.)
  2. **The Architecture:** Prototype a multi-agent system with **at least 2 agents** and **at most 5 agents** (for whatever definition of “agent” you believe makes sense in this context), that discovers which tools are available and sequences tool calls to **reliably** solve basic user-provided math problems (or if you prefer, mathy word problems). The agents can **only** use the available tools (**including the unreliable tool[s]**) , i.e. no LLM-hallucinated arithmetic should be used for user-facing answers (even if that arithmetic is correct, as is increasingly the case among frontier models).
    1. You might well choose to include a lightweight planning, reasoning, and/or task decomposition layer in your prototype — but unless you have a compelling justification, all *user-facing* outputs (and most intermediate outputs) should be structured or semistructured, not unstructured.
    2. **Don’t hesitate to ask us for an OpenAI API key or Anthropic API key.** Otherwise, we’re happy to reimburse these costs after submission (*within reason/at Spindle’s discretion*).
  3. **The Twist:** When your prototype identifies a sequence of tool calls that reliably or fairly reliably solves a certain class of math problem(s) **based on successful execution(s)** , it should learn to do something like (e.g.) **memoize** or **semantically cache that sequence of tool calls as a single, idempotent new `VIRTUAL TOOL`** (i.e. some learning behavior akin to \*\*“bundling” the tool calls into a *single* new idempotent tool, to which a *single* call can be made, which can be reliably invoked next time a math problem of the same or similar form is encountered).
  4. **The Finish Line:** Prove programmatically that your prototype works reasonably well (or at least that it could be *completed* to work reasonably well, if short on time).
    1. **Bonus points for using actual evals to show this.**
      1. **(If you’re an “evals-focused” candidate, consider reframing/approaching the entire task through the lens of an evals system instead, i.e. evals-driven development. Just tell us to judge your quality vs. emphasis vs. completion accordingly.)**
  5. **Bonus Points:**
    1. Create the math toolbox/interfaces in a non-Python language (ideally Rust, Go, or Typescript).
    2. If you decide to use a vector database anywhere, consider prototyping your own vector DB or VDB-like utility. (Not if this takes up all your time, though. It’s not the most important part.)
- **If you don’t have enough time for a project like this, or have alternate ideas, please let us know so we can find a path forward that we all feel good about!** Either way, we really look forward to seeing you through these next steps.